

# K-Truss Decomposition

ECE508 - MANYCORE PARALLEL ALGORITHMS

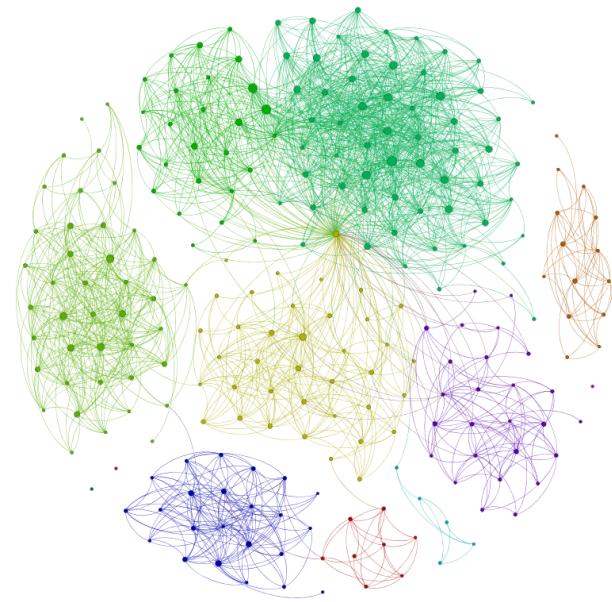
PROF. WEN-MEI HWU

HENGZHE DING, YIJIAN DUAN, DONG LIU

# Design Overview

# Problem

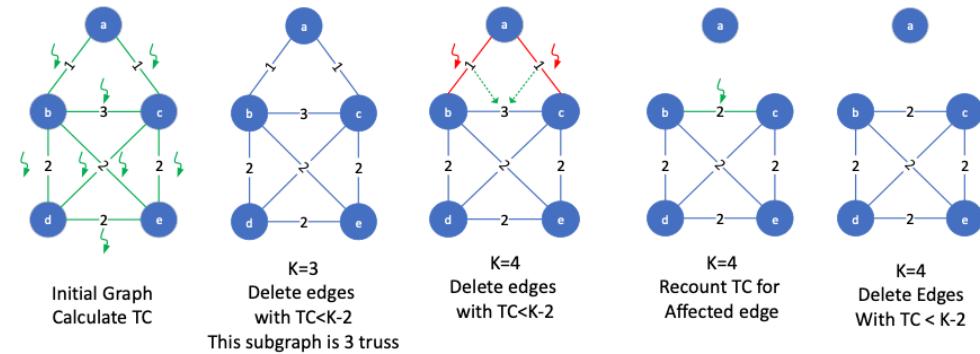
- ▶ K-truss - a type of cohesive subgraphs of a network
- ▶ High collaboration sub-network identification
- ▶ Visualization of large-scale networks and analysis of network connectivity.
- ▶ Large online social networks, e.g., Facebook and Twitter,



<http://allthingsgraphed.com/2014/08/28/facebook-friends-network/>

# Problem

- ▶ K-truss decomposition
- ▶ Measures cohesiveness of community
- ▶ Analyze community network connection



# Challenges



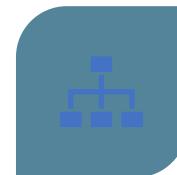
LARGE NETWORKS



BILLIONS OF EDGES  
AND MILLIONS OF  
VERTICES.



RUNNING TIME



DATA STRUCTURE



PARALLELISM

# Logic

```
While (graph is not empty for current k)
    mark all edges as affected for test tc < k-2
    While (still exist affected edges)
        test tc of affected
        if tc < k-2
            fake delete, // mark as -1
            mark affected edges // need to clear
            previous affected
        check currentEdges
```

# Implementation

# Data Structure

One thread per edge

## Graph

Shrink each round of k

- **edgeSrc**
  - # numEdges
  - Int Array
- **edgeDes**
  - # numEdges
  - Int Array
- **rowPtr**
  - # numEdges + 1
  - Int Array

## Edge flag

Inner while loop

- **affected**
  - # numEdges
  - Int Array
- **to\_delete**
  - # numEdges
  - Int Array
- **e\_aff**
  - # numEdges
  - Int Array

# Pseudocode

- Outer while loop each for one k

**Input:**  $G = (V, E)$

**Output:**  $k$ -truss for  $3 \leq k \leq k_{max}$

```

1:  $k \leftarrow 3$ 
2: Mark all  $e \in E$  as "affected"
3: for each  $e = (u, v) \in E$  do           ▷ Parallel for
4:     Binary search  $e' = (v, u) \in E$ 
5:      $\mathbf{I}_r[e] \leftarrow e'$ 
6: end for

```

- Inner while loop

```

34:  $E \leftarrow StmCmp(E, \text{"not deleted"})$       ▷ Long update
35: Output  $E$  as  $k$ -truss edge list
36: if  $E \neq \emptyset$  then  $k \leftarrow k + 1$  and goto 2

```

```

7: while true do
8:    $(E_{aff}, \mathbf{I}'_f, \mathbf{I}'_r) \leftarrow StmCmp(E, \text{"affected and } u < v\text{"})$ 
9:   if  $E_{aff} = \emptyset$  then goto 34;
10:  Mark all  $e \in E$  as "not affected"
11:  for each  $e = (u, v) \in E_{aff}$  do           ▷ Parallel for
12:     $\Delta(e) \leftarrow |\text{adj}(u) \cap \text{adj}(v)|$           ▷ Algorithm 2
13:    if  $\Delta(e) < k - 2$  then
14:      Stage 1:
15:        Lookup  $e' = (u, v) \leftarrow \mathbf{I}'_f[e]$ 
16:        Lookup  $e'' = (v, u) \leftarrow \mathbf{I}'_r[e]$ 
17:        Mark  $e'$  and  $e''$  in  $E$  as "delete"
18:      Stage 2:
19:         $W \leftarrow \text{adj}(u) \cap \text{adj}(v)$           ▷ Algorithm 2
20:        if  $W \neq \emptyset$  then
21:           $e_1 \leftarrow (u, w); e_2 \leftarrow (v, w); s.t. w \in W$ 
22:          Lookup  $e'_1 = (w, u) \leftarrow \mathbf{I}_r[e_1]$ 
23:          Lookup  $e'_2 = (w, v) \leftarrow \mathbf{I}_r[e_2]$ 
24:          Mark  $e_1, e_2, e'_1, e'_2$  as "affected"
25:        end if
26:      end if
27:    end for
28:    for each  $e = (u, v) \in E$  do           ▷ Parallel for
29:      if  $e$  labeled "delete" then
30:         $u \leftarrow -1; v \leftarrow -1$           ▷ Short update
31:      end if
32:    end for
33: end while

```

# Implementation Demo

# Parallel Optimization

```

Input:  $G = (V, E)$ 
Output:  $k$ -truss for  $3 \leq k \leq k_{max}$ 
1:  $k \leftarrow 3$ 
2: Mark all  $e \in E$  as "affected"
3: for each  $e = (u, v) \in E$  do                                ▷ Parallel for
4:   Binary search  $e' = (v, u) \in E$ 
5:    $\mathbf{I}_r[e] \leftarrow e'$ 
6: end for
7: while true do
8:    $(E_{aff}, \mathbf{I}'_f, \mathbf{I}'_r) \leftarrow StmCmp(E, "affected and u < v")$ 
9:   if  $E_{aff} = \emptyset$  then goto 34;
10:  Mark all  $e \in E$  as "not affected"
11:  for each  $e = (u, v) \in E_{aff}$  do                                ▷ Parallel for
12:     $\Delta(e) \leftarrow |\text{adj}(u) \cap \text{adj}(v)|$                                 ▷ Algorithm 2
13:    if  $\Delta(e) < k - 2$  then
14:      Stage 1:
15:      Lookup  $e' = (u, v) \leftarrow \mathbf{I}'_f[e]$ 
16:      Lookup  $e'' = (v, u) \leftarrow \mathbf{I}'_r[e]$ 
17:      Mark  $e'$  and  $e''$  in  $E$  as "delete"
18:      Stage 2:
19:       $W \leftarrow \text{adj}(u) \cap \text{adj}(v)$                                 ▷ Algorithm 2
20:      if  $W \neq \emptyset$  then
21:         $e_1 \leftarrow (u, w); e_2 \leftarrow (v, w)$ ; s.t.  $w \in W$ 
22:        Lookup  $e'_1 = (w, u) \leftarrow \mathbf{I}_r[e_1]$ 
23:        Lookup  $e'_2 = (w, v) \leftarrow \mathbf{I}_r[e_2]$ 
24:        Mark  $e_1, e_2, e'_1, e'_2$  as "affected"
25:      end if
26:    end if
27:  end for
28:  for each  $e = (u, v) \in E$  do                                ▷ Parallel for
29:    if  $e$  labeled "delete" then
30:       $u \leftarrow -1; v \leftarrow -1$                                 ▷ Short update
31:    end if
32:  end for
33: end while
34:  $E \leftarrow StmCmp(E, "not deleted")$                                 ▷ Long update
35: Output  $E$  as  $k$ -truss edge list
36: if  $E \neq \emptyset$  then  $k \leftarrow k + 1$  and goto 2

```

first mark all edge available to be affected and un-deleted

```

mark<<<dimGrid, dimBlock>>>(edgeSrcDevice, affectedDevice, e_affDevice, to_deleteDevice, numEdges);
select e_aff out of affected
selectAff<<<dimGrid, dimBlock>>>(e_aff.data(), affected.data(), currEdges, lengthDevice);
cudaDeviceSynchronize();
mark all e as "not affected"
markAll<<<dimGrid, dimBlock>>>(affectedDevice, numEdges, -1);

```

check all affected edges

```

checkAffectedEdges<<<dimGrid, dimBlock>>>(
  affected.data(), affected.size(), edgeSrcDevice, edgeDstDevice,
  to_delete.data(), e_aff.data(),
  rowPtrDevice, k, triangleCounts.data());

```

Short update

```

shortUpdate<<<dimGrid, dimBlock>>>(edgeSrcDevice, edgeDstDevice, to_delete.data(), currEdges);

```

No more long update, count current number of edges to determine k-truss availability

```

countEdges<<<dimGrid, dimBlock>>>(edgeSrcDevice, numEdges, edgeCountDevice);

```

# Results

```
int numEdges = 16;
int edgeSrc[16] = {0, 0, 1, 1, 1, 1, 2, 2, 2, 2, 3, 3, 3, 3, 4, 4, 4};
int edgeDst[16] = {1, 2, 0, 2, 3, 4, 0, 1, 3, 4, 1, 2, 4, 1, 2, 3};
int rowPtr[6] = {0, 2, 6, 10, 13, 16};
int numRow = 6;
```

```
int numEdges2 = 30;
int edgeSrc2[30] = {0, 0, 0, 1, 1, 1, 2, 2, 2, 2, 3, 3, 3, 4, 4, 4, 4, 4, 5, 5, 5, 6, 6, 6, 6, 7, 7, 7, 8, 8, 8};
int edgeDst2[30] = {1, 2, 7, 0, 3, 4, 0, 5, 6, 7, 1, 4, 8, 1, 3, 5, 8, 2, 4, 6, 2, 5, 7, 8, 0, 2, 6, 3, 4, 6};
int rowPtr2[10] = {0, 3, 6, 10, 13, 17, 20, 24, 27, 30};
int numRow2 = 10;
```

```
int numEdges3 = view.nnz();
int * edgeSrc3 = const_cast<int*>(view.row_ind());
int * edgeDst3 = const_cast<int*>(view.col_ind());
int * rowPtr3 = const_cast<int*>(view.row_ptr());
int numRow3 = view.num_rows();
```

For the first graph:

Order of truss is 3. Remaining # of edges is 16  
Order of truss is 4. Remaining # of edges is 12  
k-truss order for the first graph is: 4

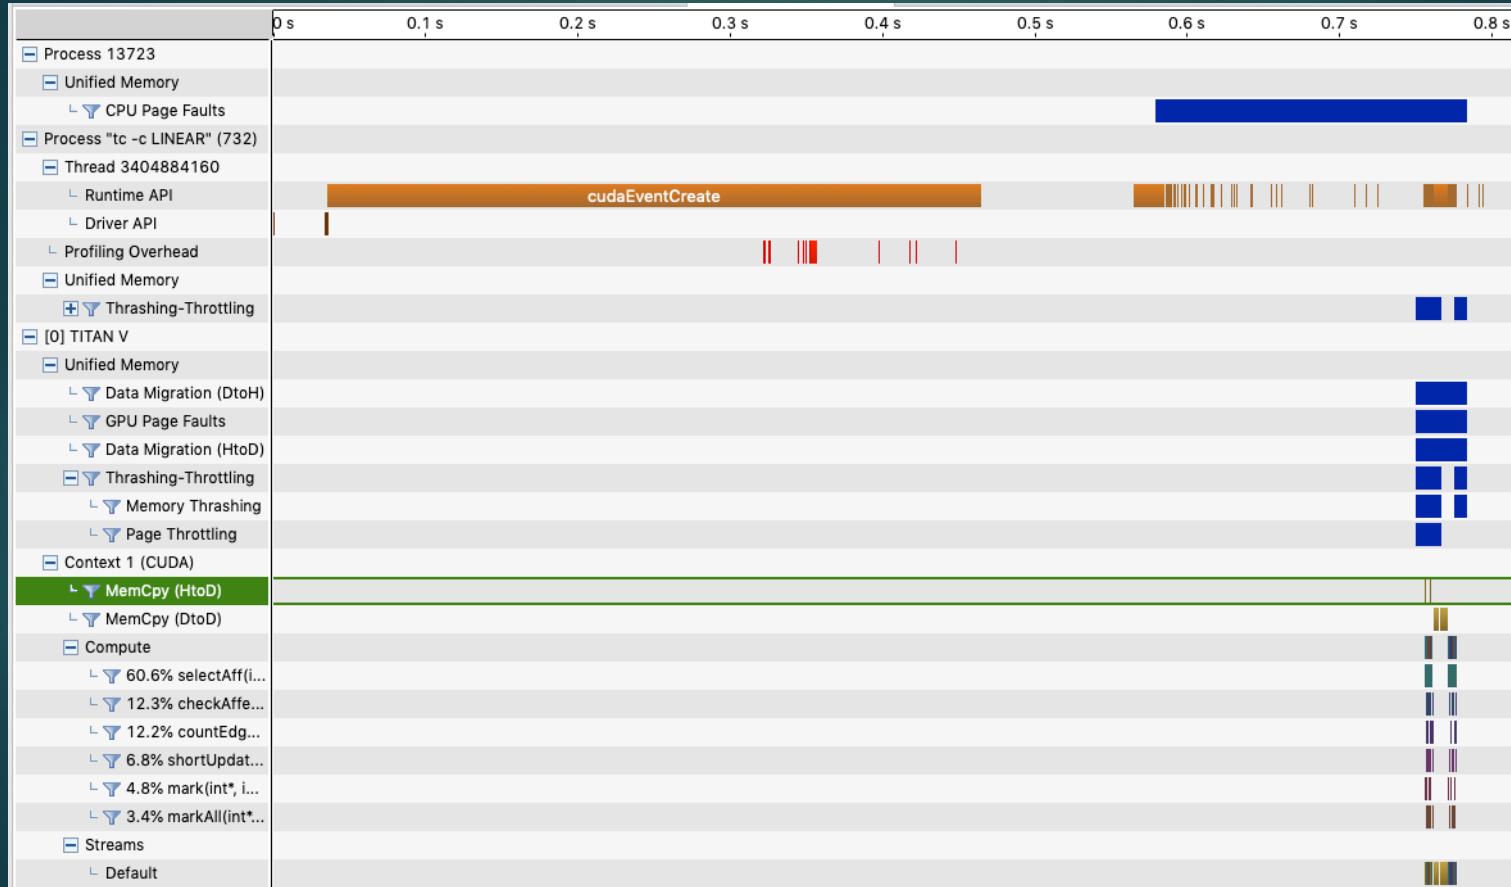
For the second graph:

Order of truss is 3. Remaining # of edges is 24  
k-truss order for the second graph is 3

For the real input graph - California Road Network:

number of edges: 5533214  
number of nodes: 1965207  
Order of truss is 3. Remaining # of edges is 719620  
Order of truss is 4. Remaining # of edges is 504  
k-truss order for the input graph is 4  
TIMER:: do k truss decomposition took 18.468191ms

# Performance analysis



# Memcpy

Memcpy (HtoD)	
▼Duration	
Session	815.56148 ms (...)
Memcpys	7.392 µs
Invocations	6
Total Bytes	432 B
Avg. Throughput	58.442 MB/s

Memcpy (DtoD)	
▼Duration	
Session	815.56148 ms (...)
Memcpys	8.4775 ms (8,4...)
Invocations	3
Total Bytes	52.127 MB
Avg. Throughput	6.149 GB/s

# GPU Details

Compute	
└	60.6% selectAff(int*, int*, int*)
└	12.3% checkAffectedEdges(int*, int*, int*)
└	12.2% countEdges(int*, int*, int*)
└	6.8% shortUpdate(int*, int*, int*)
└	4.8% mark(int*, int*, int*)
└	3.4% markAll(int*, int*, int*)

Name	Invocations	Avg. Duration	Regs	Static SMem	Avg. Dynamic SMem
markAll(int*, int, int)	14	15.556 µs	16	0	0
shortUpdate(int*, int*, int*)	14	30.811 µs	16	0	0
mark(int*, int*, int*, int*, int)	8	37.931 µs	16	0	0
checkAffectedEdges(int*, int*, int*)	14	56.038 µs	25	0	0
countEdges(int*, int, int*)	8	97.055 µs	16	0	0
selectAff(int*, int*, int, int*)	22	175.646 µs	16	0	0

# Properties

`selectAff(int*, int*, int, int*)`

▼Duration	
Session	815.56148 ms (...)
Kernel	3.86422 ms (3,...)
Invocations	22
Importance	60.6%

`countEdges(int*, int, int*)`

▼Duration	
Session	815.56148 ms (...)
Kernel	776.443 µs
Invocations	8
Importance	12.2%

`mark(int*, int*, int*, int*, int)`

▼Duration	
Session	815.56148 ms (...)
Kernel	303.454 µs
Invocations	8
Importance	4.8%

`checkAffectedEdges(int*, int, int*, int*, int*, int*, int, int)`

▼Duration	
Session	815.56148 ms (...)
Kernel	784.538 µs
Invocations	14
Importance	12.3%

`shortUpdate(int*, int*, int*, int)`

▼Duration	
Session	815.56148 ms (...)
Kernel	431.356 µs
Invocations	14
Importance	6.8%

`markAll(int*, int, int)`

▼Duration	
Session	815.56148 ms (...)
Kernel	217.789 µs
Invocations	14
Importance	3.4%

# Q & A